


# OpenCL execution model

Introduction to Vitis





# Goal



- ▶ Understand
  - What OpenCL is and how it can be used in Vitis
  - The OpenCL execution model for Xilinx accelerators



# What is OpenCL?

- ▶ OpenCL is open standard for cross-platform, parallel programming for diverse range of devices.
- ▶ Advantages
  - C/C++ based language for writing acceleration functions
  - Portable code across CPUs, GPUs and FPGAs
- ▶ Caveats
  - Performance is not portable
  - Separate optimization effort is required for each target platform

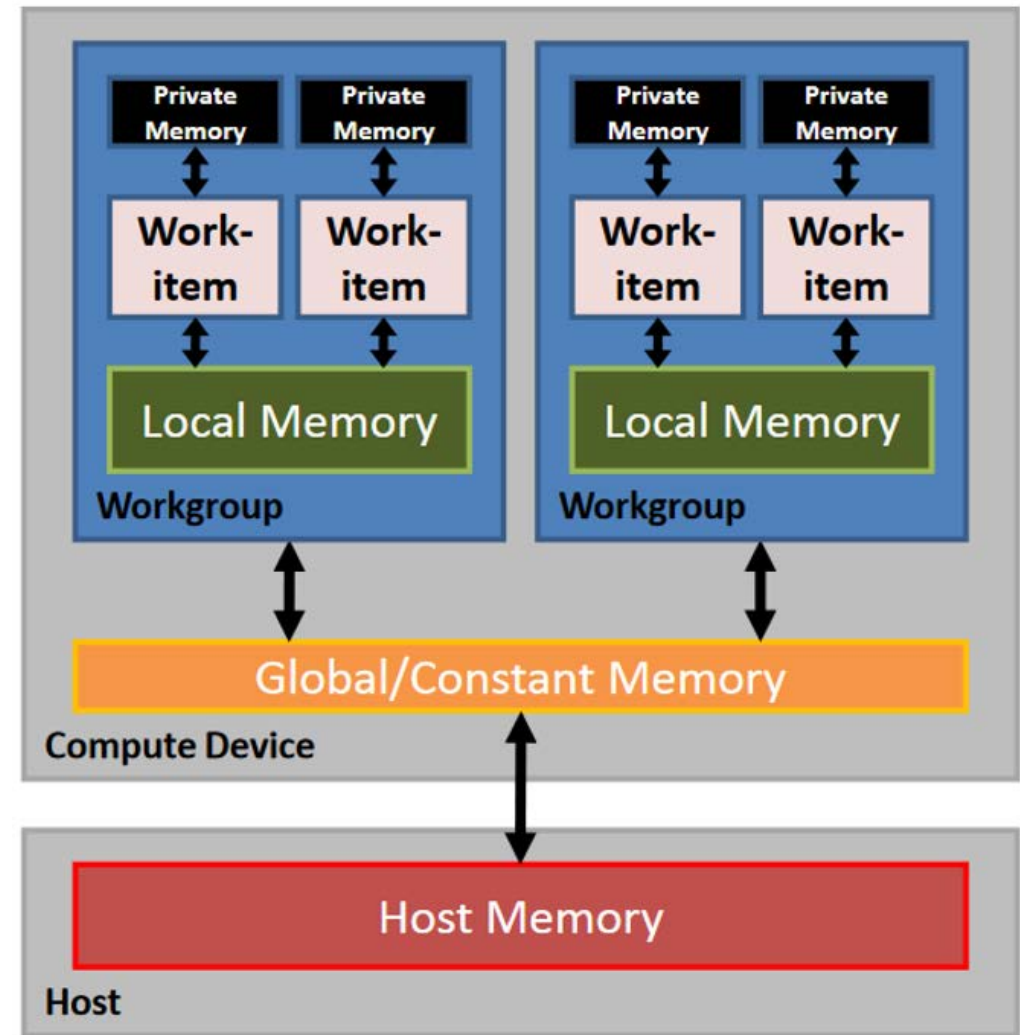


# Basic definitions

- ▶ Host code:
  - Runs on host computer. Controls devices, kernels and data transfers with the kernel
- ▶ Kernel code:
  - The code that is executed on hardware
- ▶ Platform (OpenCL):
  - In OpenCL context – a data structure that identifies vendor's OpenCL implementation. Used to access a specific device.
- ▶ Platform (Xilinx):
  - Implementation of static part on Xilinx device. Contains all necessary IPs and logic for FPGA to enable communication between host and kernel(s).

# OpenCL Memory Objects

- ▶ **Host Memory**
  - OpenCL only defines how host memory interacts with OpenCL objects
- ▶ **Global Memory**
  - Visible to host and device
  - Accessible to all work items in all workgroups
- ▶ **Constant Memory**
  - Read-only region of a global memory
- ▶ **Local Memory**
  - Local to a workgroup (shared by all work-items in a group)
- ▶ **Private Memory**
  - Accessible by a work-item





# Five structures of OpenCL application

- ▶ First step of any OpenCL application – develop host code
  - Dispatches kernels to connected devices
- ▶ Host code needs to have following structures:
  - cl\_device\_id – represents the device
  - cl\_kernel – represents specific kernel
  - cl\_program – represents the source of where kernels come from
  - cl\_command\_queue – represents the queue which passes kernels to devices
  - cl\_context – enables devices to receive kernels and transfer data




# Host Code Execution on a Heterogeneous System

- ▶ Seven steps for application execution on a heterogeneous system
  - 1. Powering up
  - 2. Initializing runtime components
  - 3. Configuring the device
  - 4. Allocating buffers
  - 5. Writing the buffers to FPGA memory
  - 6. Running the accelerators
  - 7. Reading the buffers from FPGA memory

# 1. Powering Up

- ▶ On power up, only the *shell* is initialized in the FPGA
  - This enables the PCIe interface
  - The custom logic region is now available for programming by host CPU
- ▶ The shell will manage communication with the host






## 2. Initialize Runtime Components

- ▶ Creation of OpenCL context and device
  - Device → FPGA
  - Context → Platform
- ▶ Creation of OpenCL command queues used to send commands to the FPGA

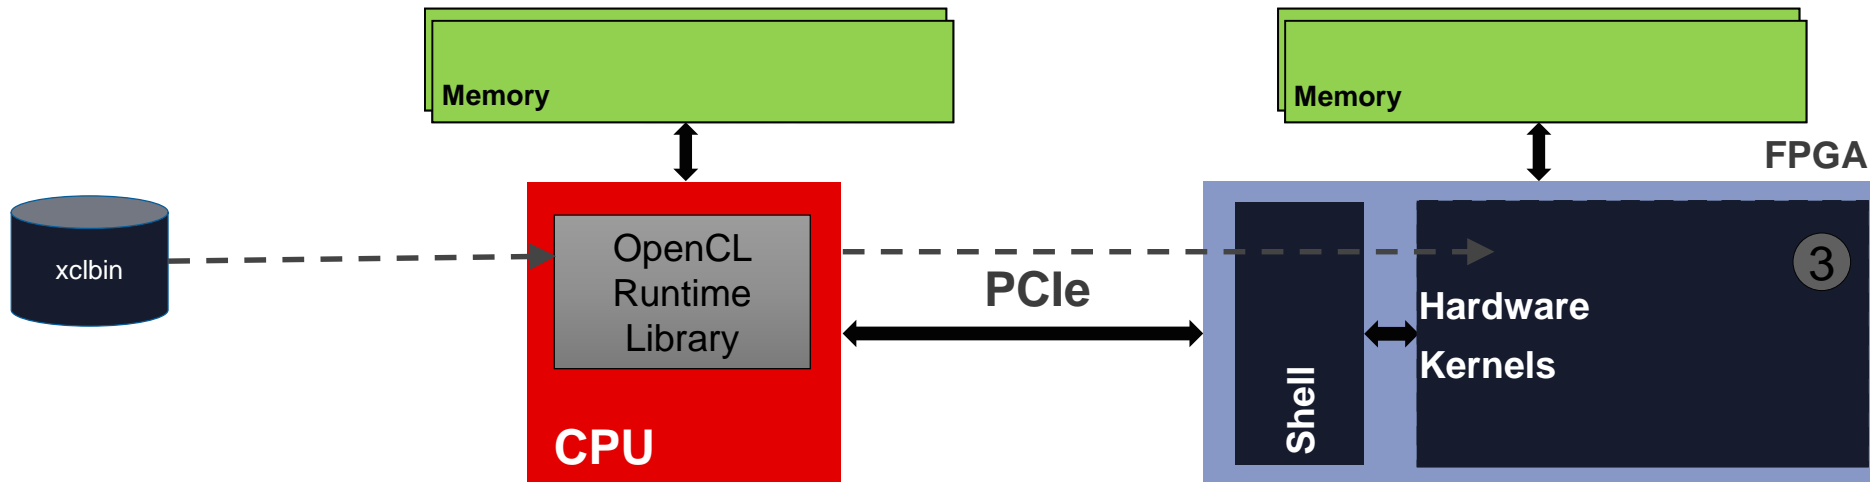
```
devices = get_devices("Xilinx");  
cl::Device device = devices[0];  
cl::Context context(...);  
  
cl::CommandQueue q(context, device,...);
```



### 3. Configuring the Device

- ▶ Create `cl::Program::Binaries` program
  - Object to store the xclbin binary file
- ▶ Host programs the FPGA by calling `cl::Program program`
  - Loads the `.xclbin` file (FPGA binary)

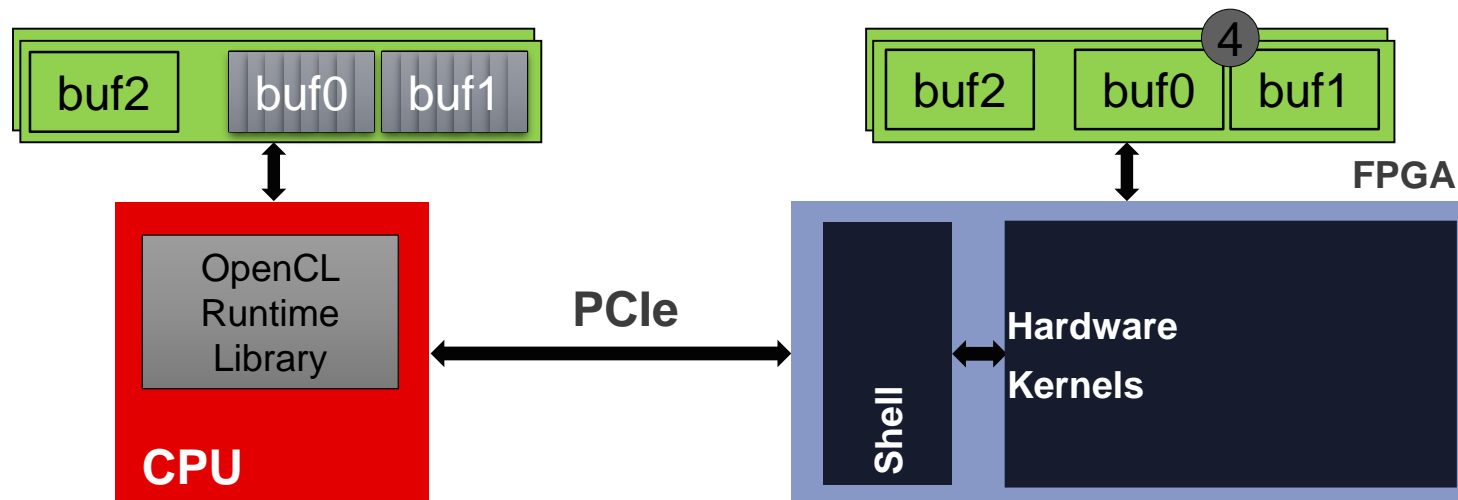
```
cl::Program::Binaries bins{{fileBuf,  
                             FileBufSize}};  
  
cl::Program program(context, devices,  
                    bins,...);
```



## 4. Allocating Buffers

- ▶ Host allocates buffers in the device
- ▶ Buffers are used to transfer data from the CPU to the FPGA and back

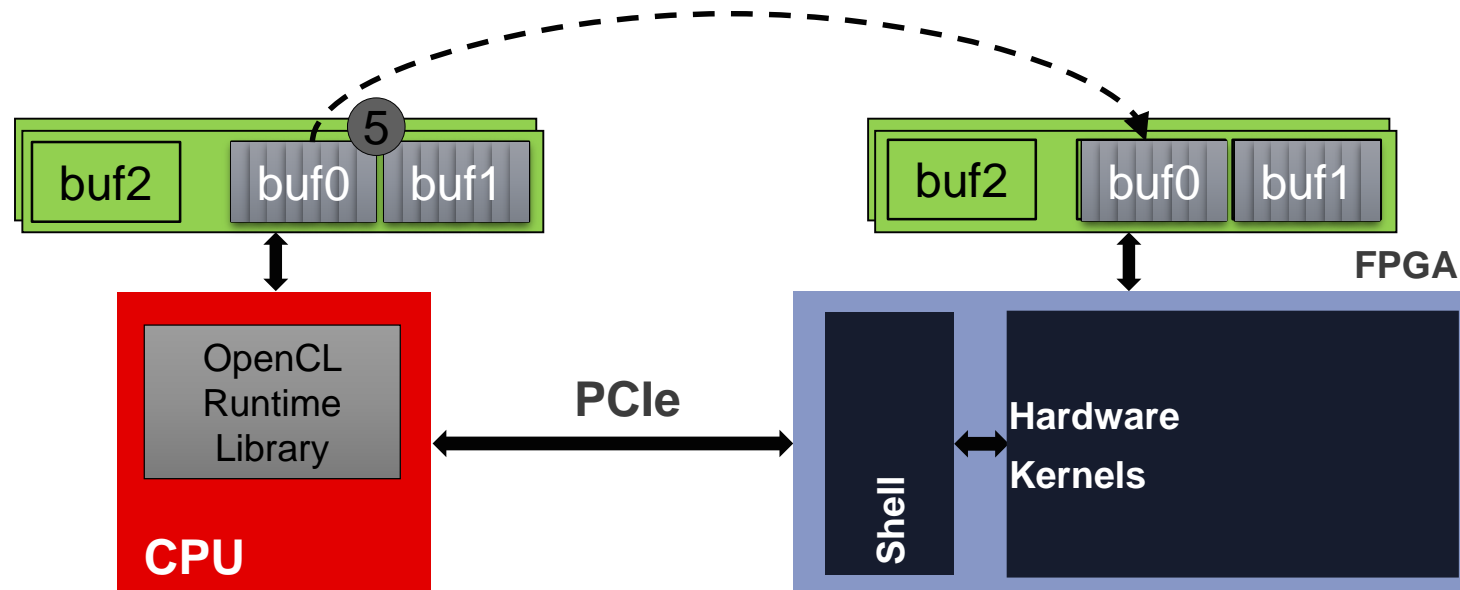
```
buf0 = cl::clCreateBuffer(context, CL_MEM_READ_ONLY,  
    ...);  
buf1 = cl::clCreateBuffer(context, CL_MEM_READ_ONLY,  
    ...);  
buf2 = cl::clCreateBuffer(context,  
    CL_MEM_WRITE_ONLY, ...);
```



## 5. Writing the Buffers to FPGA Memory

- ▶ Host copies data to be processed from host memory to the buffer in global memory

```
clEnqueueMigrateMemObjects(Command_Queue, 1,  
                             &GlobMem_BUF_DataIn_1, ..);
```




## 6. Running the Accelerators

- ▶ Host schedules execution of the desired kernel with EnqueueTask
- ▶ Runtime starts the kernel
- ▶ Kernel processes data previously copied to from host buffer to global memory

```
kernel = cl::clCreateKernel(program,
                             "mykernel", ...);

cl::clSetKernelArg(kernel, 0, sizeof(cl_mem), &buf0);
cl::clSetKernelArg(kernel, 1, sizeof(cl_mem), &buf1);
cl::clSetKernelArg(kernel, 2, sizeof(cl_mem), &buf2);


cl::clEnqueueTask(Command_Queue, Kernel, 0, NULL,
                  NULL);
```



## 6. Running the Accelerators (Cont'd)

- ▶ After the kernel finishes processing the data, it notifies the host

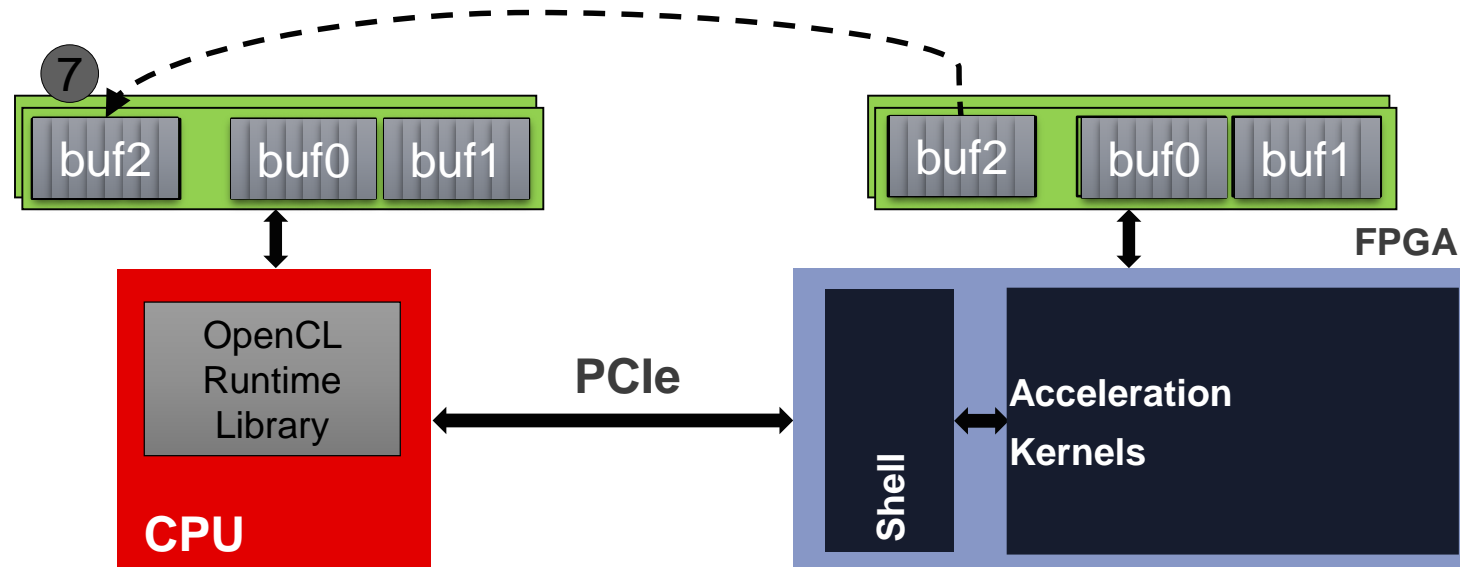
```
cl::clFinish(Command_Queue);
```



## 7. Reading the Buffers from FPGA Memory

- ▶ Host retrieves the results by scheduling a copy of global memory content back to host memory

```
cl::clEnqueueMigrateMemObjects(Command_Queue, 1, &GlobMem_BUF_RES, CL_MIGRATE_MEM_OBJECT_HOST, ...);
```





---

# Thank You

